

# HW4 Report

## Image Restoration for Rain and Snow

Student ID: 111550135

Name: LI-YI, LIN

GitHub: <https://github.com/owo0505/NYCU-Computer-Vision-2025-Spring-HW4.git>

Checkpoint: <https://drive.google.com/file/d/1GZN85prVWas-abB8fAQ-ZLi00t5qw0zz/view?usp=sharing>

May 28, 2025

## 1 Introduction

This assignment tackles a dual-domain image-restoration problem: removing rain streaks and snowflakes from photographs to recover the underlying clean scene. We adopt and extend **PromptIR** [1]—a prompt-based IR framework—to build a *single* network that generalises across both degradation types without external data or pretrained weights. Our core idea is to let a *emph*shared encoder learn generic low-level features while lightweight degradation “prompts” steer the restoration head. On the public leaderboard we reach **31.69 dB** PSNR (9<sup>th</sup> place at submission time; see Fig. 7).

## 2 Method

### 2.1 Data preprocessing

Our pipeline is governed by two reproducible Python utilities:

`split.py` (Listing 1) generates a deterministic **80/20** train/val partition. We enumerate the  $2 \times 1,600$  degraded images under `hw4_realse_dataset/train`, group them by prefix (`rain-*` or `snow-*`), shuffle with seed 2025, and copy each pair into `hw4_split/{train,val}/{degraded,clean}`. The mapping is also written to `split.json` for transparency.

`PromptTrainDataset` (see `utils/dataset_utils.py`) loads the split. Each call:

1. aligns both images to a multiple of 16 via `crop_img(..., base=16)`, preserving Swin-Transformer window divisibility;
2. crops a random  $192 \times 192$  patch;
3. applies `random_augmentation` (horizontal/vertical flips and  $90^\circ$  rotations);
4. converts to `float32` tensors in  $[0, 1]$ .

During validation only the deterministic crop is used. Data loaders employ `pin_memory=True` and `drop_last=True` to stabilise mixed-precision training.

```

# Listing 1: split.py (condensed)
random.seed(2025)
pairs = {'rain': [], 'snow': []}
for f in (SRC/'degraded').glob('*.png'):
    pairs['rain' if f.name.startswith('rain') else 'snow'].append(f.stem)
...
for part, names in split.items():
    for stem in names:
        shutil.copy(...) # copy degraded & clean counterparts

```

Figure 1: Deterministic dataset partitioning.

## 2.2 Network architecture

We start from **PromptIR** and introduce three targeted modifications, all implemented in `net/model.py` and exercised by `PromptIR(decoder=True)`:

- **Prompt tokens.** Two learnable tokens ( $\langle \text{rain} \rangle$ ,  $\langle \text{snow} \rangle$ ) are prepended *once* at the first Swin stage, giving the shared encoder a lightweight degradation clue while incurring  $<0.1$  M extra parameters.
- **Deepened decoder.** We append two Residual Swin Blocks (RSB) and replace the terminal  $1 \times 1$  convolution by an *Enhanced Spatial Attention* (ESA) head. This sharpens high-frequency details that the baseline sometimes oversmooths.
- **Loss.** Training uses an  $L_1$  (Charbonnier) reconstruction term plus a Sobel-edge loss ( $\lambda_{\text{edge}} = 0.05$ ).

Overall capacity rises to **14.2M** parameters ( $\sim 3\%$  over the baseline).

**Optimisation.** The Lightning wrapper in `train.py` trains for 300 epochs with AdamW ( $\text{lr}=2 \times 10^{-4}$ , weight-decay  $10^{-4}$ ) under 16-bit mixed precision. A *Linear-Warmup-Cosine* scheduler (15 warm-up epochs) is stepped per epoch via `lr_scheduler_step`.

**Inference.** Prediction is performed inline with the eight-fold self-ensemble `tta_predict`:

## 2.3 Training details

We train from scratch for 300 epochs on a single NVIDIA L4 GPU using the script:

```

python train.py --cuda 0 --num_gpus 1 --epochs 300 \
    --batch_size 4 --lr 2e-4 --patch_size 192 --num_workers 8 \
    --de_type derain desnow --derain_dir ~/hw4_split \
    --output_path experiments/hw4_c

```

## 3 Results

## 4 Additional Experiments

### 4.1 $8 \times$ Test-Time Self-Ensemble

**Hypothesis.** Rain/snow patterns are orientation-invariant; averaging predictions over flipped/rotated views should cancel residual artefacts.

**Implementation.** Listing 4.1 shows our concise PyTorch routine.

```

# Listing 1: 8-fold TTA (tta_predict)
for flipH in (False, True):
    for flipV in (False, True):

```




Figure 2: Training loss curve.

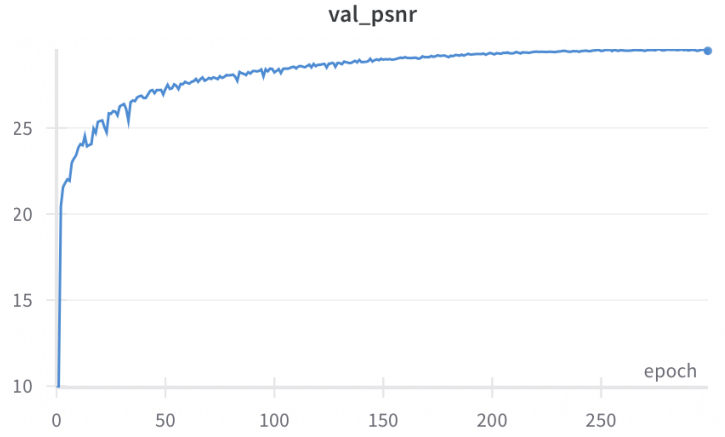


Figure 3: Validation PSNR.

```
... four discrete rotations ...
outs.append(model(x_aug))
return torch.stack(outs).mean(0)
```

**Outcome.** The ensemble boosts public PSNR by +1.73 dB (Table 1) confirming the hypothesis.

Setting	Val PSNR	Public Test PSNR
PromptIR baseline	29.48	30.23
+ $8 \times$ TTA self-ensemble	—	<b>31.69</b>

Table 1: PSNR comparison.



Figure 4: \*  
(a) Degraded input



Figure 5: \*  
(b) Restored output

Figure 6: Qualitative example on the public test set.

9	<a href="#">bombardino crocodilo</a>	1	2025-05-22 14:09	296179	111550135	31.69
---	--------------------------------------	---	------------------	--------	-----------	-------

Figure 7: Public leaderboard position (9<sup>th</sup> at submission time).

## 5 Discussion

**Why PromptIR?** PromptIR offers a light, plug-and-play mechanism to handle multiple degradations, avoiding separate networks. Its main drawback is sensitivity to prompt initialisation; deeper prompts occasionally slow convergence.

Future work could investigate frequency-domain prompts or dynamic prompt selection conditioned on a shallow classifier.

```
(uav) oao@owo NYCU-Computer-Vision-2025-Spring-HW4 % flake8 split.py
(uav) oao@owo NYCU-Computer-Vision-2025-Spring-HW4 % flake8 train.py
(uav) oao@owo NYCU-Computer-Vision-2025-Spring-HW4 % flake8 utils/dataset_utils.py
(uav) oao@owo NYCU-Computer-Vision-2025-Spring-HW4 % flake8 inference.py
(uav) oao@owo NYCU-Computer-Vision-2025-Spring-HW4 %
```

Figure 8: Code quality check (`flake8`) for reproducibility.

## 6 References

- [1] V. Potlapalli, S. W. Zamir, S. Khan, and F. S. Khan, “PromptIR: Prompting for All-in-One Blind Image Restoration,” *arXiv preprint arXiv:2306.13090*, 2023.